

Minería de datos y Patrones

Version 2024-I

HoG - Human Detection

[Capítulo 2]

Dr. José Ramón Iglesias

DSP-ASIC BUILDER GROUP

Director Semillero TRIAC

Ingeniería Electronica

Universidad Popular del Cesar

HoG: Histogram of oriented gradients

Histograms of oriented gradients for human detection



Navneet Dalal

Authors Navneet Dalal, Bill Triggs

Publication date 2005/6/20

Conference 2005 IEEE computer society conference on computer vision and pattern recognition (CVPR'05)


Volume 1

Pages 886-893

Publisher IEEE

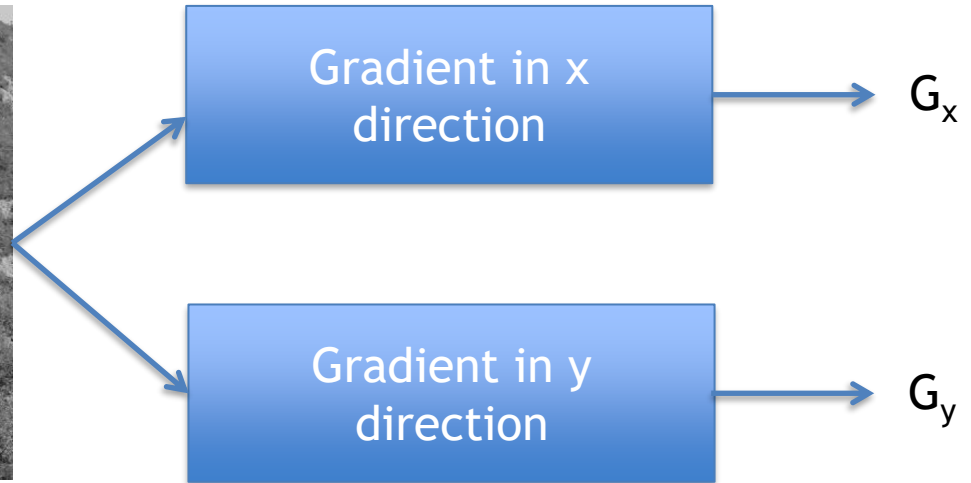
Description We study the question of feature sets for robust visual object recognition; adopting linear SVM based human detection as a test case. After reviewing existing edge and gradient based descriptors, we show experimentally that grids of histograms of oriented gradient (HOG) descriptors significantly outperform existing feature sets for human detection. We study the influence of each stage of the computation on performance, concluding that fine-scale gradients, fine orientation binning, relatively coarse spatial binning, and high-quality local contrast normalization in overlapping descriptor blocks are all important for good results. The new approach gives near-perfect separation on the original MIT pedestrian database, so we introduce a more challenging dataset containing over 1800 annotated human images with a large range of pose variations and backgrounds.

Total citations [Cited by 35511](#)




N Dalal, B Triggs: [Histograms of oriented gradients for human detection](#). Computer Vision and Pattern Recognition, (CVPR 2005).

Histogram of Gradients

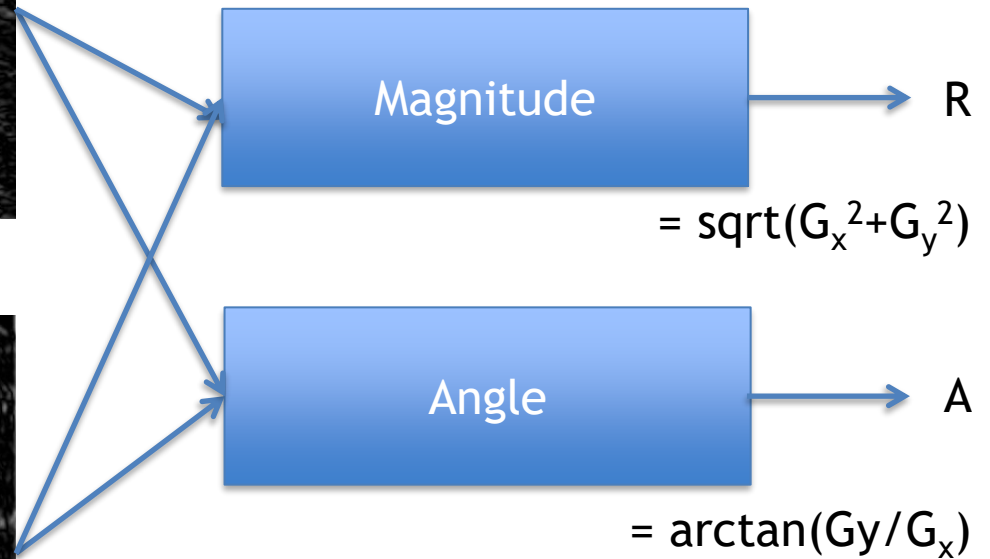



Histogram of Gradients

G_x : Gradient in x direction



G_y : Gradient in y direction




Histogram of Gradients



R: Magnitude



A: Angle



Histogram of 8 directions




Histogram of Gradients



R: Magnitude



A: Angle



Histogram of 8 directions
(computation of first bin)



Histogram of Gradients



R: Magnitude



A: Angle between -22.5° and 22.5°




Histogram of 8 directions
(computation of first bin)



Histogram of Gradients


R: Magnitude in this direction




A: Angle between -22.5° and 22.5°



Histogram of 8 directions
(computation of first bin)




Σ



Histogram of Gradients


R: Magnitude in this direction



A: Angle between 22.5° and 67.5°



Histogram of 8 directions
(computation of second bin)




Histogram of Gradients



R: Magnitude



A: Angle



Histogram of 8 directions




Histogram of Gradients



R: Magnitude





A: Angle





Histogram of 8 directions





The descriptor proposed by the authors is a concatenation of HoG in different overlapped partitions



















Descriptor of n elements

$$n = \# \text{ cells} \times \# \text{ bins}$$








Example using Python:

```
def hog(img, orientations=9,  
pixels_per_cell=(16,16),cells_per_block=(2,2),norm=False):  
    X = skimage.feature.hog(img, orientations=orientations,  
pixels_per_cell=pixels_per_cell,cells_per_block=cells_per_block)  
    if norm:  
        X = X/np.linalg.norm(X)  
    return X
```

La función `hog()` toma 6 parámetros como entrada:



- **image**: la imagen de destino a la que desea aplicar la extracción de características HOG.
- **Orientaciones**: Número de contenedores en el histograma que queremos crear, el trabajo de investigación original utilizó 9 contenedores, por lo que pasaremos 9 como orientaciones.
- **pixels_per_cell**: Determina el tamaño de la celda, como mencionamos anteriormente, es 8x8.
- **cells_per_block**: El número de celdas por bloque, será de 2x2 como se mencionó anteriormente.
- **visualizar**: Un booleano si para devolver la imagen del HOG, la establecemos en **True** para que podamos mostrar la imagen.
- **multicanal**: Lo establecemos en **True** para indicar a la función que la última dimensión se considera como un canal de color, en lugar de espacial.

Example using Balu:



How to detect pedestrians using HoG?

The Solution




Strategy: Sliding Windows



Strategy: Sliding Windows



Strategy: Sliding Windows



The Key-Idea:

Design a classifier that is able to distinguish between pedestrians from no-pedestrians.

Positive Class: Pedestrians




Pedestrian

Negative Class: No-Pedestrians




No-Pedestrian

Construction of the Dataset: Positive Class




Original Image : 720 x 475

Crop



100 x 200

Resize




64 x 128

Construction of the Dataset: Negative Class



Construction of the Dataset: Negative Class



Negative Class: No-Pedestrians



00000265a.png



00000266a.png



00000267a.png



00000272a.png



00000274a.png



00000275a.png



00000276a.png



00000277a.png



00000278a.png



00000280a.png



00000281a.png



00000282a.png



00000283a.png



00000284a.png



00000285a.png



00000286a.png



00000287a.png



00000288a.png



00000289a.png



00000290a.png



00000292a.png



00000293a.png



00000295a.png



00000299a.png



00000300a.png



00000303a.png



00000305a.png



00000306a.png



00000307a.png



00000308a.png



00000309a.png



00000310a.png



00000311a.png



00000312a.png



00000313a.png



00000314a.png



00000315a.png



00000316a.png



00000317a.png



00000318a.png



00000319a.png



00000320a.png



00000321a.png



00000322a.png



00000323a.png



00000324a.png



00000325a.png



00000326a.png



00000327a.png



00000328a.png



00000329a.png



00000330a.png



00000331a.png



00000332a.png



00000333a.png



00000334a.png



00000335a.png



00000336a.png



00000338a.png



00000340a.png



00000341a.png



00000342a.png



00000343a.png



00000344a.png



00000345a.png



00000365a.png



00000369a.png



00000378a.png



00000382a.png



00000383a.png



00000400a.png



00000410a.png



00000414a.png



00000415a.png



00000416a.png



00000420a.png



00000421a.png



00000422a.png



00000423a.png



00000424a.png



00000425a.png



00000426a.png



00000427a.png






00000428a.png

Image:
128 x 64 x 3 pixels

Feature Extraction

Descriptor
3780 elements



Center : The RGB patch and gradients represented using arrows. Right : The gradients in the same patch represented as numbers

Dataset: Positive Class



Pedestrian	Element-1	Element-2	Element-3	Element-4	Element-5	...	Element-3780
1	0.3452	0.1151	0.2685	0.1342	0.2416		0.2301
2	0.2301	0.1534	0.2148	0.1611	0.2071		0.3068
3	0.1726	0.1918	0.1790	0.1879	0.1812		0.3836
4	0.1381	0.2301	0.1534	0.2148	0.1611		0.4603
5	0.1151	0.2685	0.1342	0.2416	0.1450		0.5370
6	0.0986	0.3068	0.1193	0.2685	0.1318		0.6137
7	0.0863	0.3452	0.1074	0.2953	0.1208		0.6904
8	0.0767	0.3836	0.0976	0.3222	0.1115		0.7671
9	0.0690	0.4219	0.0895	0.3490	0.1036		0.8438
10	0.0628	0.4603	0.0826	0.3759	0.0967		0.9205
11	0.0575	0.4986	0.0767	0.4027	0.0906		0.9972
12	0.0531	0.5370	0.0716	0.4296	0.0853		1.0740
13	0.0493	0.5753	0.0671	0.4564	0.0805		1.1507
14	0.0460	0.6137	0.0632	0.4833	0.0763		1.2274
15	0.0432	0.6520	0.0597	0.5101	0.0725		1.3041
16	0.0406	0.6904	0.0565	0.5370	0.0690		1.3808
17	0.0384	0.7288	0.0537	0.5638	0.0659		1.4575
18	0.0363	0.7671	0.0511	0.5907	0.0630		1.5342
19	0.0345	0.8055	0.0488	0.6175	0.0604		1.6109
20	0.0329	0.8438	0.0467	0.6444	0.0580		1.6876
21	0.0314	0.8822	0.0447	0.6712	0.4851		1.7644
22	0.0300	0.9205	0.0430	0.0802	0.6444		1.8411
23	0.0288	0.9589	0.3593	0.0604	0.4660		1.9178
24	0.0276	0.1146	0.4773	0.0835	0.3625		0.2293
25	0.2310	0.0863	0.3452	0.1074	0.2953		0.1726
:	:	:	:	:	:		:
:	:	:	:	:	:		:
100000	0.1230	0.0493	0.5753	0.0671	0.4564		0.0986

Dataset: Negative Class



No-Pedestrian	Element-1	Element-2	Element-3	Element-4	Element-5	...	Element-3780
1	0.7452	0.2484	0.5796	0.2898	0.5216		0.4968
2	0.4968	0.3312	0.4637	0.3478	0.4471		0.6624
3	0.3726	0.4140	0.3864	0.4057	0.3912		0.8280
4	0.2981	0.4968	0.3312	0.4637	0.3478		0.9936
5	0.2484	0.5796	0.2898	0.5216	0.3130		1.1592
6	0.2129	0.6624	0.2576	0.5796	0.2845		1.3248
7	0.1863	0.7452	0.2318	0.6376	0.2608		1.4904
8	0.1656	0.8280	0.2108	0.6955	0.2408		1.6560
9	0.1490	0.9108	0.1932	0.7535	0.2236		1.8216
10	0.1355	0.9936	0.1783	0.8114	0.2087		1.9872
11	0.1242	1.0764	0.1656	0.8694	0.1956		2.1528
12	0.1146	1.1592	0.1546	0.9274	0.1841		2.3184
13	0.1065	1.2420	0.1449	0.9853	0.1739		2.4840
14	0.0994	1.3248	0.1364	1.0433	0.1647		2.6496
15	0.0932	1.4076	0.1288	1.1012	0.1565		2.8152
16	0.0877	1.4904	0.1220	1.1592	0.1490		2.9808
17	0.0828	1.5732	0.1159	1.2172	0.1423		3.1464
18	0.0784	1.6560	0.1104	1.2751	0.1361		3.3120
19	0.0745	1.7388	0.1054	1.3331	0.1304		3.4776
20	0.0710	1.8216	0.1008	1.3910	0.1252		3.6432
21	0.0677	1.9044	0.0966	1.4490	0.4851		3.8088
22	0.0648	1.9872	0.0927	0.3740	1.3910		3.9744
23	0.0621	2.0700	0.3593	0.1304	1.0060		4.1400
24	0.0596	0.5342	1.0304	0.1803	0.7825		1.0684
25	0.2310	0.1863	0.7452	0.2318	0.6376		0.3726
:	:	:	:	:	:		:
:	:	:	:	:	:		:
800000	0.9230	0.1065	1.2420	0.1449	0.9853		0.2129

Training





Image:
128 x 64 x 3 pixels

Descriptor
3780 elements



Output:
0 / 1
Pedestrian

Feature Extraction
Using HoG Features

Classification
Using Trained Model





Example



Example



Example





Detection

Example



Example





Detection

Detection

Detection

Example